

0. Importing Libraries and Files

```
import pandas as pd
import geopandas as geo_pd
import numpy as np
import seaborn as sns
from imblearn.over_sampling import SMOTE
import csv
import ast
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split, GridSearchCV,
RandomizedSearchCV, validation_curve, cross_validate, cross_val_score
from sklearn.preprocessing import label_binarize, LabelEncoder
from sklearn.metrics import accuracy_score, classification_report
from sklearn.metrics import roc_curve, auc
from itertools import cycle
import matplotlib.pyplot as plt
from matplotlib.legend_handler import HandlerLine2D
from scipy.stats import randint
import xgboost as xgb

def read_data(path):
    data = pd.read_csv(path, skipinitialspace=True)
    return data

# Part A
location_data = read_data('./all_data/partA/data/location_2021.csv')
train_cases = read_data('./all_data/partA/data/cases_2021_train.csv')
world_data =
geo_pd.read_file('./world_map_data/ne_110m_admin_0_countries.shp')
# Part B
processed_train_cases =
read_data('./all_data/partB/data/cases_2021_train_processed_2.csv')
processed_test_cases =
read_data('./all_data/partB/data/cases_2021_test_processed_unlabelled_
2.csv')
```

1. Data Visualization

Data Reading and Cleaning

```
def clean_cases(cases : pd.DataFrame):
    ''' Various values found from train cases
    'Hospitalized' 'Recovered' 'Deceased' 'Discharged' 'Alive'
    'Discharge'
    'Under treatment' 'Stable' 'Died' 'Receiving treatment' 'Death'
    'Stable condition' 'Dead' 'Discharged from hospital' 'Critical
```

```

condition'
    'Released from quarantine' 'Recovering at home 03.03.2020'
    '''
    hospitalized = ['Hospitalized', 'Discharged', 'Discharge', 'Under
treatment',
                    'Receiving treatment', 'Discharged from hospital',
    'Critical condition']
    non_hospitalized = ['Alive', 'Stable', 'Stable condition',
    'Released from quarantine',
                        'Recovering at home 03.03.2020', 'Recovered']
    deceased = ['Deceased', 'Died', 'Death', 'Dead']

    cases.loc['outcome'] = cases['outcome'].str.capitalize()
    cases.loc[cases['outcome'].isin(hospitalized), 'outcome'] =
    'Hospitalized'
    cases.loc[cases['outcome'].isin(non_hospitalized), 'outcome'] =
    'Non-Hospitalized'
    cases.loc[cases['outcome'].isin(deceased), 'outcome'] = 'Deceased'
    # Handle cases where there is more cases than what is known. Can't
    use 'Other' due to project req.
    cases.loc[~cases['outcome'].isin(['Hospitalized', 'Non-
Hospitalized', 'Deceased']), 'outcome'] = 'Non-Hospitalized'
    return

def clean_dataset(location_data: pd.DataFrame, train_cases:
pd.DataFrame):
    location_data = location_data[['Country_Region', 'Confirmed',
    'Deaths', 'Province_State']]
    #
    age,sex,province,country,latitude,longitude,date_confirmation,addition
    al_information,source,chronic_disease_binary
    train_cases = train_cases[['age', 'sex', 'country', 'province',
    'chronic_disease_binary', 'outcome', 'outcome_group']]
    # standardize naming convention
    clean_cases(train_cases)
    location_data = location_data.rename(columns={'Country_Region':
    'country', 'Province_State': 'province'})
    location_data.loc[location_data['country'] == 'US', 'country'] =
    'United States'

    # Fix missing data
    location_data['province'] =
    location_data['province'].fillna(location_data['country'])

    return train_cases, location_data

# Colour Palette with Indices for Seaborn
# Reference: https://stackoverflow.com/questions/36271302/changing-color-scale-in-seaborn-bar-plot
def colors_from_values(values, palette_name):

```

```

    # normalize the values to range [0, 1]
    normalized = (values - min(values)) / (max(values) - min(values))
    # convert to indices
    indices = np.round(normalized * (len(values) -
1)).astype(np.int32)
    # use the indices to get the colors
    palette = sns.color_palette(palette_name, len(values))
    return list(np.array(palette).take(indices, axis=0))

```

Data Plotting

```

def plot_data(df: pd.DataFrame, world: geo_pd.GeoDataFrame):
    total_available_data = df['count'].sum()

    merged_world = pd.merge(world, df, left_on=['NAME_CIAWF'],
right_on=['country'], how='left')

    # Percentage is based on available data of country / total
    # available data from all countries
    merged_world['available_percentage'] = (merged_world['count'] /
total_available_data) * 100
    merged_world['available_percentage_log'] =
np.log(merged_world['available_percentage'])

    # Plot Heat Map
    _, ax = plt.subplots(1, 1)
    merged_world.plot(column='available_percentage', ax=ax,
cmap="rocket_r",
                    legend=True, legend_kwds={'shrink': 0.6},
vmax=100, vmin=0) # Only used for proper Legend values
    merged_world.plot(column='available_percentage_log', ax=ax,
cmap="rocket_r", legend=False) # Using Log for bigger differences
    ax.tick_params(left=False, labelleft=False, bottom=False,
labelbottom=False)

    # # Plot Bar Plot
    by_continent =
merged_world.groupby('CONTINENT').sum('count').reset_index()
    by_continent.loc[by_continent['CONTINENT'] == 'Seven seas (open
ocean)', 'CONTINENT'] = 'Seven seas'
    by_continent['log_count'] =
np.log(by_continent['count'].replace(0.0, np.nan))
    by_continent.fillna(0, inplace=True)

    y_values_log = by_continent['log_count'].values.tolist()

    _, ax1 = plt.subplots(1, 1)
    ax1 = sns.barplot(by_continent, x = 'CONTINENT', y = 'log_count',
                    palette =
colors_from_values(np.array(y_values_log), "YlOrRd"), hue =

```

```

'CONTINENT', legend = False)
ax1.set(xlabel='Continents', ylabel='Available Data')

# Add value to the bars
for p in ax1.patches:
    # Revert the log to use correct numbers when showcasing values
    value = np.round(np.exp(p.get_height()),
decimals=0).astype(int) if p.get_height() > 0 else 0
    ax1.annotate(value,
                  xy = (p.get_x()+p.get_width()/2., p.get_height()),
                  ha = 'center',
                  va = 'center',
                  xytext = (0, 10),
                  textcoords = 'offset points')

plt.tick_params(left = False, labelleft = False)
plt.xticks(rotation = 45)
plt.show()

```

Driver Code

```

train_data, location_data = clean_dataset(location_data, train_cases)


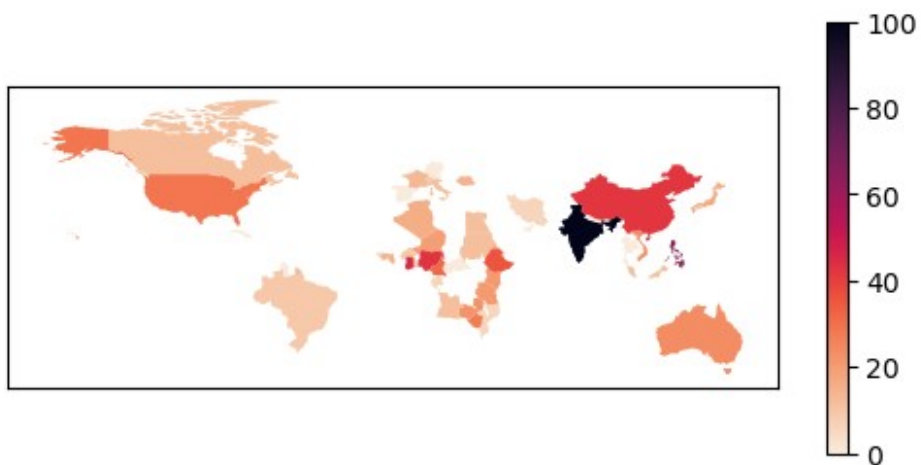
# Merge data set
merged_data = pd.concat([train_cases, location_data], ignore_index =
True) # concat on country
num_available_data_by_country =
merged_data[merged_data['outcome'].notna()] \
            .value_counts('country').reset_index()

# Plot
plot_data(num_available_data_by_country, world_data)

C:\Users\pewpe\AppData\Local\Temp\ipykernel_78132\832819255.py:14:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#
returning-a-view-versus-a-copy
cases.loc['outcome'] = cases['outcome'].str.capitalize()

```



2. Data Pre-processing

Preprocessing

```
# Merge data
merged_data = pd.merge(train_cases, location_data, how='left',
on=['country', 'province'])
merged_data = merged_data.dropna(subset =
merged_data.columns.difference(['outcome_group']))
# Add Expected Mortality Rate
# Observed deaths in cases_2021_train (based on location) / Location
(country, province) Deaths count
observed_deaths_data = merged_data[merged_data['outcome'] ==
'Deceased'] \
    .value_counts(['country', 'province',
'outcome']).reset_index() \
    .rename(columns={"count":
"observed_deaths"}).drop('outcome', axis=1)
merged_data = pd.merge(merged_data, observed_deaths_data, how='left',
on=['country', 'province'])
merged_data['Expected_Mortality_Rate'] =
merged_data['observed_deaths'] / merged_data['Deaths']
```

Driver Code

merged_data

	age	sex	province	country	latitude	longitude	\
0	26	male	Bihar	India	26.221520	84.358790	
1	42	female	Queensland	Australia	-28.016700	153.400000	
2	30	male	Karnataka	India	12.867673	75.251507	
3	4	male	Karnataka	India	13.335010	74.749360	
4	50	male	Karnataka	India	12.867673	75.251507	
...	
3891	21	male	Karnataka	India	17.900000	77.550000	
3892	25	female	Karnataka	India	16.200000	77.366670	
3893	28	male	Karnataka	India	15.433120	75.635450	
3894	38	male	Shaanxi	China	34.955530	109.858500	
3895	39	male	Bihar	India	25.752980	86.027930	

	date_confirmation	additional_information	\
0	07.04.2020		Same family,
1	30.01.2020	QLD Case 2 quarantine at the Gold Coast	Univer...
2	26.05.2020		International Travel from Qatar
3	21.05.2020		Travelled from Mumbai, Maharashtra

4	28.05.2020	Returnee from\r\
nMaharashtra		
...	...	
...		
3891	31.05.2020	Travelled from
Maharashtra		
3892	31.05.2020	Travelled from
Maharashtra		
3893	14.05.2020	Travelled from Ahmedabad,
Gujarat		
3894	03.02.2020	son-in-law of another
patient		
3895	08.05.2020	from outside the state and are in quarantine
f...		

	source \
0	https://twitter.com/sanjayjavin/status/1247571...
1	https://www.health.qld.gov.au/news-events/doh-...
2	https://twitter.com/DHFWKA/status/126518463555...
3	https://twitter.com/DHFWKA/status/126337507668...
4	https://t.me/Karnataka_KoViD19_Broadcast/3670
...	...
3891	https://twitter.com/DHFWKA/status/126707739902...
3892	https://twitter.com/DHFWKA/status/126707739902...
3893	https://twitter.com/DHFWKA/status/126082976107...
3894	http://sxwjw.shaanxi.gov.cn/art/2020/2/3/art_9...
3895	https://twitter.com/sanjayjavin/status/1258594...

	chronic_disease_binary	outcome	outcome_group	Confirmed Deaths \
0	False	Hospitalized	NaN	265527.0
1576.0				
1	False	discharge	NaN	1477.0
6.0				
2	False	Hospitalized	NaN	997004.0
12567.0				
3	False	Hospitalized	NaN	997004.0
12567.0				
4	False	Hospitalized	NaN	997004.0
12567.0				
...
...				
3891	False	Hospitalized	NaN	997004.0
12567.0				
3892	False	Hospitalized	NaN	997004.0
12567.0				
3893	False	Hospitalized	NaN	997004.0
12567.0				
3894	False	stable	NaN	567.0

```

3.0
3895                False  Hospitalized                NaN    265527.0
1576.0

```

	observed_deaths	Expected_Mortality_Rate
0	3.0	0.001904
1	NaN	NaN
2	23.0	0.001830
3	23.0	0.001830
4	23.0	0.001830
...
3891	23.0	0.001830
3892	23.0	0.001830
3893	23.0	0.001830
3894	NaN	NaN
3895	3.0	0.001904

```
[3896 rows x 16 columns]
```

3. Feature Selection

Feature Removing

```

test_cases = processed_test_cases
train_cases = processed_train_cases

# Combine Country and Province and remove their individual columns
train_cases['combined_key'] = train_cases['country'] + ', ' +
train_cases['province']
test_cases['combined_key'] = test_cases['country'] + ', ' +
test_cases['province']
train_cases.drop('country', axis=1, inplace=True)
test_cases.drop('country', axis=1, inplace=True)
train_cases.drop('province', axis=1, inplace=True)
test_cases.drop('province', axis=1, inplace=True)

# These features do not seem to impact the performance
train_cases.drop('Confirmed', axis=1, inplace=True)
test_cases.drop('Confirmed', axis=1, inplace=True)
train_cases.drop('Active', axis=1, inplace=True)
test_cases.drop('Active', axis=1, inplace=True)
train_cases.drop('Incident_Rate', axis=1, inplace=True)
test_cases.drop('Incident_Rate', axis=1, inplace=True)
train_cases.drop('Recovered', axis=1, inplace=True)
test_cases.drop('Recovered', axis=1, inplace=True)
train_cases.drop('Deaths', axis=1, inplace=True)
test_cases.drop('Deaths', axis=1, inplace=True)

```


Driver Code

```
print(train_cases.head(10))
print(test_cases.head(10))
```

	age	sex	latitude	longitude	date_confirmation	\
0	18	female	25.490960	85.939030	2020-05-18	
1	27	female	7.070000	125.600000	2020-04-15	
2	46	male	13.083620	80.282520	2020-05-02	
3	21	female	13.083620	80.282520	2020-05-24	
4	27	male	26.283610	87.203470	2020-05-27	
5	24	male	24.457120	85.137490	2020-05-26	
6	65	female	19.420820	76.050130	2020-04-03	
7	35	male	25.473982	84.536523	2020-04-29	
8	22	male	19.387650	85.050120	2020-05-02	
9	36	female	20.189990	86.304550	2020-04-30	

	chronic_disease_binary	Case_Fatality_Ratio	outcome_group	\
0	False	0.593537	hospitalized	
1	False	1.779368	nonhospitalized	
2	False	1.434463	hospitalized	
3	False	1.434463	hospitalized	
4	False	0.593537	hospitalized	
5	False	0.593537	hospitalized	
6	False	1.942744	deceased	
7	False	0.593537	hospitalized	
8	False	0.563480	hospitalized	
9	False	0.563480	hospitalized	

	combined_key
0	India, Bihar
1	NaN
2	India, Tamil Nadu
3	India, Tamil Nadu
4	India, Bihar
5	India, Bihar
6	India, Maharashtra
7	India, Bihar
8	India, Odisha
9	India, Odisha

	age	sex	latitude	longitude	date_confirmation
0	59	female	14.59580	120.97720	2020-03-31
1	79	male	11.13927	79.08428	2020-05-24
2	44	female	13.08362	80.28252	2020-05-19
3	36	male	13.12462	79.91815	2020-04-30
4	52	male	25.31258	86.48888	2020-04-24

False					
5	28	female	12.68224	79.98008	2020-05-31
False					
6	17	male	25.73271	86.98845	2020-05-23
False					
7	18	male	13.00287	76.10245	2020-05-24
False					
8	68	male	14.63000	121.03000	2020-04-11
False					
9	21	male	13.08362	80.28252	2020-05-15
False					

	Case_Fatality_Ratio	combined_key
0	1.779368	NaN
1	1.434463	India, Tamil Nadu
2	1.434463	India, Tamil Nadu
3	1.434463	India, Tamil Nadu
4	0.593537	India, Bihar
5	1.434463	India, Tamil Nadu
6	0.593537	India, Bihar
7	1.260476	India, Karnataka
8	1.779368	NaN
9	1.434463	India, Tamil Nadu

4. Mapping the Features

Feature Mapping

```
# 0: deceased, 1: hospitalized, 2: non_hospitalized.
# unique() == ['hospitalized' 'nonhospitalized' 'deceased']
train_cases['outcome_group'].replace(to_replace =
train_cases['outcome_group'].unique(),
                                     value=[1, 2, 0], inplace=True)

# 0: Female, 1: Male
# unique() == ['female' 'male']
train_cases['sex'].replace(to_replace = train_cases['sex'].unique(),
value=[0, 1], inplace=True)
test_cases['sex'].replace(to_replace = test_cases['sex'].unique(),
value=[0, 1], inplace=True)

# 0: False, 1: True
# unique() == [False True]
train_cases['chronic_disease_binary'].replace(to_replace =
train_cases['chronic_disease_binary'].unique(),
                                               value=[0, 1], inplace=True)
test_cases['chronic_disease_binary'].replace(to_replace =
test_cases['chronic_disease_binary'].unique(),
                                              value=[0, 1],
inplace=True)
le = LabelEncoder()
```

```

train_cases['combined_key'] =
le.fit_transform(train_cases['combined_key'])
test_cases['combined_key'] =
le.fit_transform(test_cases['combined_key'])

# Parse the date_confirmation into date_time format
train_cases['date_confirmation'] =
pd.to_datetime(train_cases['date_confirmation'], errors='coerce')
test_cases['date_confirmation'] =
pd.to_datetime(test_cases['date_confirmation'], errors='coerce')

# Extracting year, month, and day
train_cases['year'] = train_cases['date_confirmation'].dt.year
train_cases['month'] = train_cases['date_confirmation'].dt.month
train_cases['day'] = train_cases['date_confirmation'].dt.day
train_cases['dayofweek'] =
train_cases['date_confirmation'].dt.dayofweek

test_cases['year'] = test_cases['date_confirmation'].dt.year
test_cases['month'] = test_cases['date_confirmation'].dt.month
test_cases['day'] = test_cases['date_confirmation'].dt.day
test_cases['dayofweek'] = test_cases['date_confirmation'].dt.dayofweek

# drop unnecessary dates
test_cases.drop('date_confirmation', axis = 1, inplace = True)
train_cases.drop('date_confirmation', axis = 1, inplace = True)

```

Driver Code

```

print(train_cases.head(10))
print(test_cases.head(10))

```

	age	sex	latitude	longitude	chronic_disease_binary	\
0	18	0	25.490960	85.939030	0	
1	27	0	7.070000	125.600000	0	
2	46	1	13.083620	80.282520	0	
3	21	0	13.083620	80.282520	0	
4	27	1	26.283610	87.203470	0	
5	24	1	24.457120	85.137490	0	
6	65	0	19.420820	76.050130	0	
7	35	1	25.473982	84.536523	0	
8	22	1	19.387650	85.050120	0	
9	36	0	20.189990	86.304550	0	

	Case_Fatality_Ratio	outcome_group	combined_key	year	month	day
\						
0	0.593537	1	13	2020	5	18

1	1.779368	2	44	2020	4	15
2	1.434463	1	35	2020	5	2
3	1.434463	1	35	2020	5	24
4	0.593537	1	13	2020	5	27
5	0.593537	1	13	2020	5	26
6	1.942744	0	27	2020	4	3
7	0.593537	1	13	2020	4	29
8	0.563480	1	31	2020	5	2
9	0.563480	1	31	2020	4	30

dayofweek

0	0
1	2
2	5
3	6
4	2
5	1
6	4
7	2
8	5
9	3

age	sex	latitude	longitude	chronic_disease_binary
Case_Fatality_Ratio \				
0 59	0	14.59580	120.97720	0
1.779368				
1 79	1	11.13927	79.08428	0
1.434463				
2 44	0	13.08362	80.28252	0
1.434463				
3 36	1	13.12462	79.91815	0
1.434463				
4 52	1	25.31258	86.48888	0
0.593537				
5 28	0	12.68224	79.98008	0
1.434463				
6 17	1	25.73271	86.98845	0
0.593537				
7 18	1	13.00287	76.10245	0
1.260476				
8 68	1	14.63000	121.03000	0
1.779368				

9	21	1	13.08362	80.28252	0
1.434463					

	combined_key	year	month	day	dayofweek
0	35	2020	3	31	1
1	28	2020	5	24	6
2	28	2020	5	19	1
3	28	2020	4	30	3
4	10	2020	4	24	4
5	28	2020	5	31	6
6	10	2020	5	23	5
7	20	2020	5	24	6
8	35	2020	4	11	5
9	28	2020	5	15	4

5. Balancing the Classes

Removing Outliers

```
def reduce_outliers(df, column_names, remove=False):
    outlier_indices = []
    for column in column_names:
        Q1 = df[column].quantile(0.25)
        Q3 = df[column].quantile(0.75)
        IQR = Q3 - Q1
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR

        column_outliers = df.index[(df[column] < lower_bound) |
(df[column] > upper_bound)]
        outlier_indices.extend(column_outliers)

    outlier_indices = list(set(outlier_indices)) # Remove duplicates
    if remove:
        return df.drop(index=outlier_indices)
    else:
        return df.loc[outlier_indices]

def visualize_outliers(df, column_names):
    for column in column_names:
        plt.figure(figsize=(10, 6))
        df.boxplot(column=[column])
        plt.title(f"Box plot for {column}")
        plt.show()
```

Driver Code

```
column_names_to_check = ['Case_Fatality_Ratio']

# Visualizing outliers before removal
print("Before outlier removal:")
visualize_outliers(train_cases, column_names_to_check)

# Removing outliers
train_cases_no_outliers = reduce_outliers(train_cases,
column_names_to_check, remove=True)
# Visualizing data after outlier removal
print("After outlier removal:")
visualize_outliers(train_cases_no_outliers, column_names_to_check)
''' We discover that removing the outliers make the prediction worst
'''

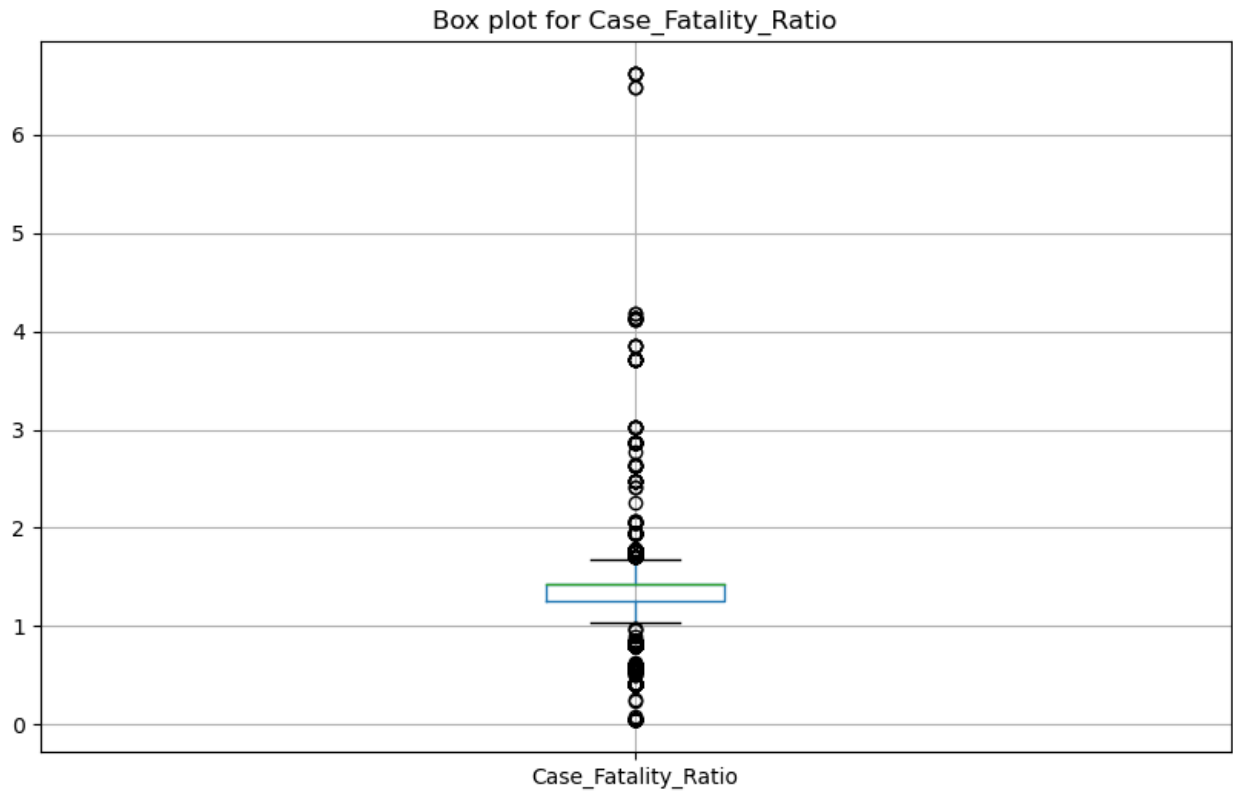
# For train_cases
# Display original outcome distribution in the training set
train_cases_outcome = train_cases['outcome_group'].value_counts()
print("Original outcome distribution in training data:\n",
train_cases_outcome)

# Separating features and target variable in the training dataset
X_train = train_cases.drop('outcome_group', axis=1)
y_train = train_cases['outcome_group']

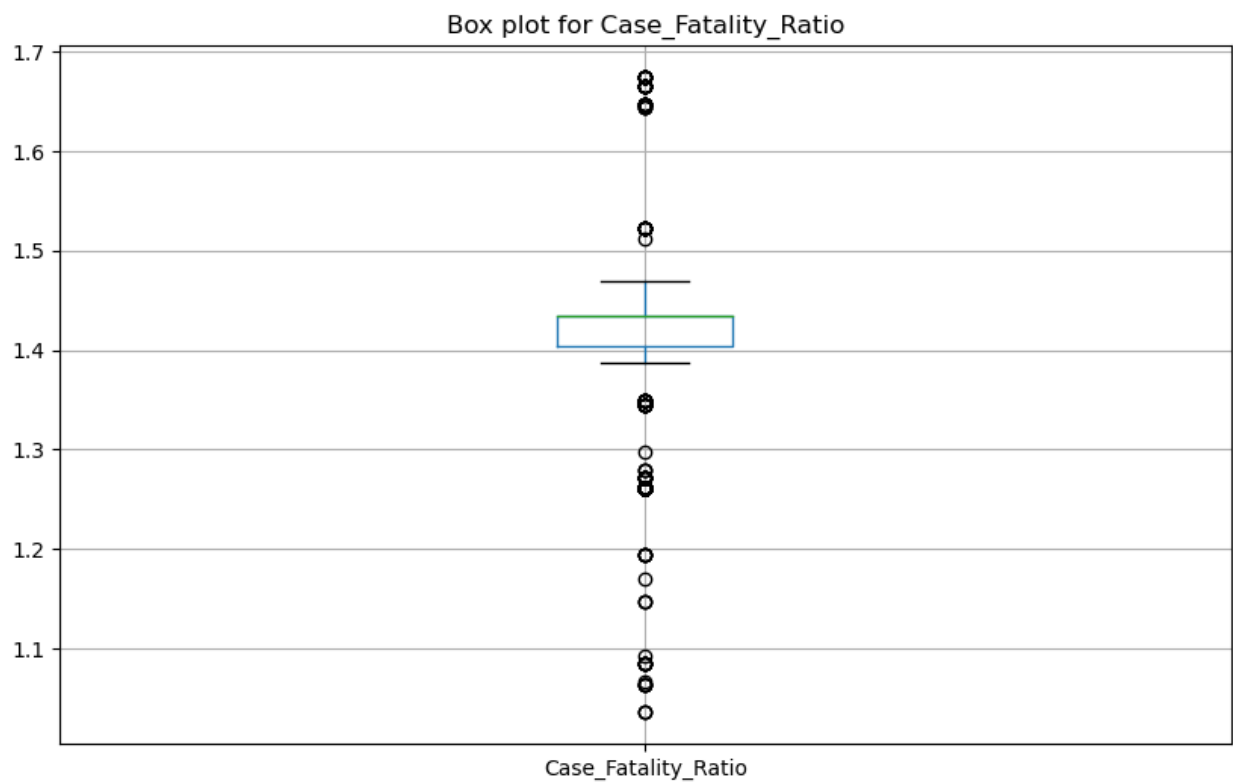
# Applying SMOTE to balance the classes in the training dataset
# random_state can be change between 0~42 (doesn't matter)
smote = SMOTE(random_state=42)
X_train_resampled, y_train_resampled = smote.fit_resample(X_train,
y_train)

balanced_train_cases = pd.concat([X_train_resampled,
y_train_resampled], axis=1)
print("Resampled outcome distribution in training data:\n",
      balanced_train_cases['outcome_group'].value_counts())
train_cases = balanced_train_cases

Before outlier removal:
```



After outlier removal:



```

Original outcome distribution in training data:
outcome_group
1      13241
2       2974
0        997
Name: count, dtype: int64
Resampled outcome distribution in training data:
outcome_group
1      13241
2      13241
0      13241
Name: count, dtype: int64

```

6. Building Models and Hyperparameter Tuning

AUC-ROC Functions

```

# https://stackoverflow.com/questions/51378105/plot-multi-class-roc-curve-for-decisiontreeclassifier
# Credits to drew_psy
# Calculating False Positive / True Positive rates
def calc_roc_auc(clf, x, y):
    y_proba = clf.predict_proba(x)
    y_bin = label_binarize(y, classes = [0, 1, 2])
    n_classes = y_bin.shape[1]

    fpr = dict()
    tpr = dict()
    roc_auc = dict()
    for i in range(n_classes):
        fpr[i], tpr[i], _ = roc_curve(y_bin[:, i], y_proba[:, i])
        roc_auc[i] = auc(fpr[i], tpr[i])
    return fpr, tpr, roc_auc, n_classes

def plot_roc_model(clf, x, y):
    fpr, tpr, roc_auc, n_classes = calc_roc_auc(clf, x, y)

    colors = cycle(['blue', 'red', 'green'])
    for i, color in zip(range(n_classes), colors):
        plt.plot(fpr[i], tpr[i], color=color,
                 label='ROC curve of class {0} (area = {1:0.2f})'
                 ''.format(i, roc_auc[i]))

    plt.plot([0, 1], [0, 1], 'k--')
    plt.xlim([-0.05, 1.0])
    plt.ylim([0.0, 1.05])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('Receiver operating characteristic for multi-class

```



```

data')
plt.legend(loc="lower right")
plt.show()
return

```

Model Definitions

```

RANDOM_STATE = 42
TARGET_NAMES = ['deceased', 'hospitalized', 'non_hospitalized']

def model_rf(x_train, x_test, y_train, y_test, params: dict,
             hyper_tuning: bool = False):
    if not hyper_tuning:
        clf = RandomForestClassifier(**params)
        clf.fit(x_train, y_train)

        train_preds = clf.predict(x_train)
        train_acc = accuracy_score(y_train, train_preds)
        y_preds = clf.predict(x_test)
        test_acc = accuracy_score(y_test, y_preds)

        rep = classification_report(y_test, y_preds, zero_division =
1, output_dict=True, target_names=TARGET_NAMES)
        results = {
            'hyperparameters': str(params),
            'mean_macro_f1': f"{rep['macro avg']['f1-score']: .2f}",
            'mean_deceased_f1': f"{rep['deceased']['f1-score']: .2f}",
            'mean_overall_accuracy': f"{test_acc: .2f}"
        }

        # plot_roc_model(clf, x_test, y_test)
        _, _, train_roc_auc, _ = calc_roc_auc(clf, x_train, y_train)
        _, _, test_roc_auc, _ = calc_roc_auc(clf, x_test, y_test)
        avg_train_auc = np.mean(list(train_roc_auc.values()))
        avg_test_auc = np.mean(list(test_roc_auc.values()))
        return train_acc, test_acc, rep, results, avg_train_auc,
avg_test_auc
    else:
        clf = RandomForestClassifier()
        tuning = RandomizedSearchCV(clf, param_distributions=params,
random_state=RANDOM_STATE, n_iter=50, cv=5, verbose=1,
return_train_score=True)
        tuning.fit(x_train, y_train)
        results = pd.DataFrame(tuning.cv_results_)
        results.to_csv('./all_data/partB/model_results/%s.csv' % 'rf')

        train_preds = tuning.predict(x_train)
        train_acc = accuracy_score(y_train, train_preds)
        y_preds = tuning.predict(x_test)

```

```

        test_acc = accuracy_score(y_test, y_preds)
        rep = classification_report(y_test, y_preds, zero_division =
1, output_dict=True, target_names=TARGET_NAMES)
        # best_score: average cross-validated score
        return train_acc, test_acc, rep, tuning.best_score_,
tuning.best_params_

def model_xgboost(x_train, x_test, y_train, y_test, params: dict,
hyper_tuning: bool = False):
    if not hyper_tuning:
        clf = xgb.XGBClassifier(objective="multi:softprob",
random_state=RANDOM_STATE, **params)          # for multi-class
classification
        clf.fit(x_train, y_train)

        train_preds = clf.predict(x_train)
        train_acc = accuracy_score(y_train, train_preds)
        y_preds = clf.predict(x_test)
        test_acc = accuracy_score(y_test, y_preds)

        rep = classification_report(y_test, y_preds, zero_division =
1, output_dict=True, target_names=TARGET_NAMES)
        results = {
            'hyperparameters': str(params),
            'mean_macro_f1': f"{rep['macro avg']['f1-score']: .2f}",
            'mean_deceased_f1': f"{rep['deceased']['f1-score']: .2f}",
            'mean_overall_accuracy': f"{test_acc: .2f}"
        }

        return train_acc, test_acc, rep, results, y_preds
    else:
        clf = xgb.XGBClassifier(objective="multi:softprob",
random_state=RANDOM_STATE)
        tuning = RandomizedSearchCV(clf, param_distributions=params,
random_state=RANDOM_STATE, n_iter=50, cv=5, verbose=1, n_jobs=1,
return_train_score=True, scoring='accuracy')
        tuning.fit(x_train, y_train)
        results = pd.DataFrame(tuning.cv_results_)
        results.to_csv('./all_data/partB/model_results/%s.csv' %
'xgb')

        train_preds = tuning.predict(x_train)
        train_acc = accuracy_score(y_train, train_preds)
        y_preds = tuning.predict(x_test)
        test_acc = accuracy_score(y_test, y_preds)
        rep = classification_report(y_test, y_preds, zero_division =
1, output_dict=True, target_names=TARGET_NAMES)
        return train_acc, test_acc, rep, tuning.best_score_,

```

```

tuning.best_params_

def model_knn(x_train, x_test, y_train, y_test, params: dict,
hyper_tuning: bool = False):

    if not hyper_tuning:
        clf = KNeighborsClassifier(**params)
        clf.fit(x_train, y_train)

        train_preds = clf.predict(x_train)
        train_acc = accuracy_score(y_train, train_preds)
        y_preds = clf.predict(x_test)
        test_acc = accuracy_score(y_test, y_preds)

        rep = classification_report(y_test, y_preds, zero_division =
1, output_dict=True, target_names=TARGET_NAMES)
        results = {
            'hyperparameters': str(params),
            'mean_macro_f1': f"{rep['macro avg']['f1-score']: .2f}",
            'mean_deceased_f1': f"{rep['deceased']['f1-score']: .2f}",
            'mean_overall_accuracy': f"{test_acc: .2f}"
        }
        return train_acc, test_acc, rep, results

    else:
        clf = KNeighborsClassifier()
        # tuning = RandomizedSearchCV(clf, param_distributions=params,
random_state=RANDOM_STATE,
                                # n_iter=10, cv=5, verbose=2,
n_jobs=1, return_train_score=True)
        tuning = GridSearchCV(clf, param_grid=params, refit=True,
verbose=3, n_jobs=-1)

        tuning.fit(x_train, y_train)
        results = pd.DataFrame(tuning.cv_results_)
        results.to_csv('./all_data/partB/model_results/%s.csv' %
'knn')

        train_preds = tuning.predict(x_train)
        train_acc = accuracy_score(y_train, train_preds)
        y_preds = tuning.predict(x_test)
        test_acc = accuracy_score(y_test, y_preds)
        rep = classification_report(y_test, y_preds, zero_division =
1, output_dict=True, target_names=TARGET_NAMES)
        # best_score: average cross-validated score
        return train_acc, test_acc, rep, tuning.best_score_,
tuning.best_params_

```

Driver Code

```
X = train_cases.drop('outcome_group', axis=1).values
y = train_cases['outcome_group'].values

X_test_cases = test_cases.values

# Do a 80/20 train test split
x_train, x_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=RANDOM_STATE)
```

Random Forest

```
# Smaller max_features reduces overfitting; sqrt is best for
classification generally
# https://datascience.stackexchange.com/questions/66825/how-many-
features-does-random-forest-need-for-the-trees
rf_params = {'n_estimators': 40, 'criterion': 'gini',
             'min_samples_split': 2, 'min_samples_leaf': 1,
             'max_features': 'sqrt', 'max_depth': 20,
             'random_state': RANDOM_STATE
            }

rf_params_tuning = {'n_estimators': range(20, 60),
                   'criterion': ['gini', 'entropy'],
                   'min_samples_split': [2,3,4,5],
                   'min_samples_leaf': [1,2,3,4,5],
                   'max_features': ['sqrt', 'log2'],
                   'max_depth': range(20, 26),
                   'random_state': [RANDOM_STATE]
                  }

print("\n1. Random Forest -- Without Scalers(): ")
train_acc, test_acc, rep, results, _, _ = model_rf(x_train, x_test,
y_train, y_test, rf_params)
print('Random Forest Train Accuracy: %.2f' % train_acc)
print('Random Forest Test Accuracy: %.2f' % test_acc)
print('Report: \n', rep)
print('Results: \n')
print('Hyperparamters: %s' % results['hyperparameters'])
print('Mean macro f1-score: %s' % results['mean_macro_f1'])
print('Mean deceased f1-score: %s' % results['mean_deceased_f1'])
print('Mean overall accuracy: %s' % results['mean_overall_accuracy'])

print("\n2. Without Scalers(), with Hyperparam tuning(): ")
train_acc, test_acc, rep, best_score, best_params = model_rf(x_train,
x_test, y_train, y_test, rf_params_tuning, hyper_tuning=True)
print("Random Forest best params: ", best_params)
print('Random Forest Train Accuracy: %.2f' % train_acc)
print('Random Forest Test Accuracy: %.2f' % test_acc)
print('Random Forest Report: \n', rep)
```

```

1. Random Forest -- Without Scalers():
Random Forest Train Accuracy: 0.99
Random Forest Test Accuracy: 0.95
Report:
{'deceased': {'precision': 0.9186530457813091, 'recall':
0.9256576439191765, 'f1-score': 0.9221420432966198, 'support':
2623.0}, 'hospitalized': {'precision': 0.9632925472747497, 'recall':
0.9752252252252253, 'f1-score': 0.9692221600447678, 'support':
2664.0}, 'non_hospitalized': {'precision': 0.9573896353166986,
'recall': 0.9382994732881866, 'f1-score': 0.9477484324529737,
'support': 2658.0}, 'accuracy': 0.946507237256136, 'macro avg':
{'precision': 0.9464450761242524, 'recall': 0.946394114144196, 'f1-
score': 0.9463708785981204, 'support': 7945.0}, 'weighted avg':
{'precision': 0.9465802310504836, 'recall': 0.946507237256136, 'f1-
score': 0.9464948706590685, 'support': 7945.0}}
Results:

Hyperparamters: {'n_estimators': 40, 'criterion': 'gini',
'min_samples_split': 2, 'min_samples_leaf': 1, 'max_features': 'sqrt',
'max_depth': 20, 'random_state': 42}
Mean macro f1-score: 0.95
Mean deceased f1-score: 0.92
Mean overall accuracy: 0.95

2. Without Scalers(), with Hyperparam tuning():
Fitting 5 folds for each of 50 candidates, totalling 250 fits
Random Forest best params: {'random_state': 42, 'n_estimators': 58,
'min_samples_split': 2, 'min_samples_leaf': 1, 'max_features': 'log2',
'max_depth': 25, 'criterion': 'gini'}
Random Forest Train Accuracy: 0.99
Random Forest Test Accuracy: 0.95
Random Forest Report:
{'deceased': {'precision': 0.9279484262419416, 'recall':
0.9329012581014106, 'f1-score': 0.9304182509505703, 'support':
2623.0}, 'hospitalized': {'precision': 0.9671764267064528, 'recall':
0.9733483483483484, 'f1-score': 0.9702525724976613, 'support':
2664.0}, 'non_hospitalized': {'precision': 0.9607917776931861,
'recall': 0.9495861550037622, 'f1-score': 0.9551561021759697,
'support': 2658.0}, 'accuracy': 0.9520453115166772, 'macro avg':
{'precision': 0.9519722102138601, 'recall': 0.9519452538178403, 'f1-
score': 0.9519423085414004, 'support': 7945.0}, 'weighted avg':
{'precision': 0.9520895239631331, 'recall': 0.9520453115166772, 'f1-
score': 0.952050955942208, 'support': 7945.0}}

```

XGBoost

```

xgb_params = {"colsample_bytree": 0.89,
              "gamma": 0.07,
              "learning_rate": 0.12,

```

```

        "max_depth": 9,
        "n_estimators": 370,
        "subsample": 0.78,
    }
# For a good general understanding on main params for XGBoost:
# https://medium.com/@rithpansanga/the-main-parameters-in-xgboost-and-their-effects-on-model-performance-4f9833cac7c
xgb_params_tuning = {
    "colsample_bytree": np.linspace(0.5, 1, 100,
endpoint=True), # controls fraction of features used for each tree.

#smaller -> smaller and less complex models (prevents overfitting)
common=[0.5, 1]
    "gamma": np.linspace(0, 0.1, 100, endpoint=True),
    "learning_rate": np.linspace(0.05, 0.3, 100,
endpoint=True), # smaller -> slower but more accurate, default=0.3
    "max_depth": randint(2, 10),
# smaller -> simpler model (underfitting), larger -> overfitting.
default=6
    "n_estimators": randint(100, 500),
# number of trees. larger --> overfitting. default 100. common=[100,
1000]
    "subsample": np.linspace(0.7, 1, 100,
endpoint=True)
}

print("\n1. XGBoost -- Without Scalers(): ")
train_acc, test_acc, rep, results, y_preds = model_xgboost(x_train,
x_test, y_train, y_test, xgb_params)
print('XGBoost Train Accuracy: %.2f' % train_acc)
print('XGBoost Test Accuracy: %.2f' % test_acc)
print('XGBoost Report: \n', rep)
print('Results: \n')
print('Hyperparamters: %s' % results['hyperparameters'])
print('Mean macro f1-score: %s' % results['mean_macro_f1'])
print('Mean deceased f1-score: %s' % results['mean_deceased_f1'])
print('Mean overall accuracy: %s' % results['mean_overall_accuracy'])

print("\n2. Without Scalers(), with Hyperparam tuning(): ")
train_acc, test_acc, rep, best_score, best_params =
model_xgboost(x_train, x_test, y_train, y_test, xgb_params_tuning,
hyper_tuning=True)
print("XGBoost best params: ", best_params)
print("XGBoost best score: ", best_score)
print('XGBoost Train Accuracy: %.2f' % train_acc)
print('XGBoost Test Accuracy: %.2f' % test_acc)
print('XGBoost Report: \n', rep)

```

1. XGBoost -- Without Scalers():

```
XGBoost Train Accuracy: 0.99
XGBoost Test Accuracy: 0.95
XGBoost Report:
{'deceased': {'precision': 0.93111279333838, 'recall':
0.9378574151734655, 'f1-score': 0.9344729344729346, 'support':
2623.0}, 'hospitalized': {'precision': 0.970193740685544, 'recall':
0.9774774774774775, 'f1-score': 0.9738219895287958, 'support':
2664.0}, 'non_hospitalized': {'precision': 0.9629629629629629,
'recall': 0.9488337095560572, 'f1-score': 0.9558461246920599,
'support': 2658.0}, 'accuracy': 0.9548143486469478, 'macro avg':
{'precision': 0.954756498995629, 'recall': 0.9547228674023334, 'f1-
score': 0.9547136828979301, 'support': 7945.0}, 'weighted avg':
{'precision': 0.9548723143698447, 'recall': 0.9548143486469478, 'f1-
score': 0.9548172796297942, 'support': 7945.0}}
Results:
```

```
Hyperparamters: {'colsample_bytree': 0.89, 'gamma': 0.07,
'learning_rate': 0.12, 'max_depth': 9, 'n_estimators': 370,
'subsample': 0.78}
Mean macro f1-score: 0.95
Mean deceased f1-score: 0.93
Mean overall accuracy: 0.95
```

```
2. Without Scalers(), with Hyperparam tuning():
Fitting 5 folds for each of 50 candidates, totalling 250 fits
XGBoost best params: {'colsample_bytree': 0.8383838383838385,
'gamma': 0.05454545454545454, 'learning_rate': 0.2368686868686869,
'max_depth': 9, 'n_estimators': 222, 'subsample': 0.7484848484848484}
XGBoost best score: 0.9479828076483733
XGBoost Train Accuracy: 0.99
XGBoost Test Accuracy: 0.95
XGBoost Report:
{'deceased': {'precision': 0.9289493575207861, 'recall':
0.9370949294700724, 'f1-score': 0.9330043651546783, 'support':
2623.0}, 'hospitalized': {'precision': 0.9718256949661909, 'recall':
0.9710960960960962, 'f1-score': 0.9714607585429967, 'support':
2664.0}, 'non_hospitalized': {'precision': 0.9598028062191885,
'recall': 0.9522197140707299, 'f1-score': 0.9559962228517469,
'support': 2658.0}, 'accuracy': 0.9535556954059157, 'macro avg':
{'precision': 0.9535259529020551, 'recall': 0.953470246545633, 'f1-
score': 0.953487115516474, 'support': 7945.0}, 'weighted avg':
{'precision': 0.9536480396598562, 'recall': 0.9535556954059157, 'f1-
score': 0.9535909214473515, 'support': 7945.0}}
```

K-Nearest Neighbours

```
knn_params = {'n_neighbors': 4,
               'weights': 'distance', # Distance is is a lot more
strict
               'algorithm': 'kd_tree',
```



```

        'leaf_size': 25,
        'p': 1, # p = 1: manhattan_distance (l1), p = 2:
euclidean_distance (l2)
        'n_jobs': -1, # use all your cpu cores
    }

knn_params_tuning = {'n_neighbors': range(3, 10),
                    'weights': ['uniform', 'distance'], # Distance is
is a lot more strict
                    'algorithm': ['kd_tree', 'ball_tree'], # brute
force is worse in every aspect
                    'leaf_size': [25], # Has no effect on the model,
only on the speed of execution
                    'p': [1, 2],
                    'n_jobs': [-1],
                    }

print("\n1. KNN -- Without Scalers(): ")
train_acc, test_acc, rep, results = model_knn(x_train, x_test,
y_train, y_test, knn_params)
print('K-Nearest Train Accuracy: %.2f' % train_acc)
print('K-Nearest Test Accuracy: %.2f' % test_acc)
print('Report: \n', rep)
print('Results: \n')
print('Hyperparamters: %s' % results['hyperparameters'])
print('Mean macro f1-score: %s' % results['mean_macro_f1'])
print('Mean deceased f1-score: %s' % results['mean_deceased_f1'])
print('Mean overall accuracy: %s' % results['mean_overall_accuracy'])

print("\n2. Without Scalers(), with Hyperparam tuning(): ")
train_acc, test_acc, rep, best_score, best_params = model_knn(x_train,
x_test, y_train, y_test, knn_params_tuning, hyper_tuning=True)
print("K-Nearest best params: ", best_params)
print("K-Nearest best score: ", best_score)
print('K-Nearest Train Accuracy: %.2f' % train_acc)
print('K-Nearest Test Accuracy: %.2f' % test_acc)
print('K-Nearest Report: \n', rep)

```

```

1. KNN -- Without Scalers():
K-Nearest Train Accuracy: 0.99
K-Nearest Test Accuracy: 0.94

```

```
Report:
```

```

{'deceased': {'precision': 0.9013878743608473, 'recall':
0.9409073579870377, 'f1-score': 0.9207237455698564, 'support':
2623.0}, 'hospitalized': {'precision': 0.9739096573208723, 'recall':
0.9388138138138138, 'f1-score': 0.956039755351682, 'support': 2664.0},
'non_hospitalized': {'precision': 0.9594543387646836, 'recall':
0.9525959367945824, 'f1-score': 0.9560128374551634, 'support':
2658.0}, 'accuracy': 0.944115796098175, 'macro avg': {'precision':

```



```

0.9449172901488011, 'recall': 0.9441057028651446, 'f1-score':
0.9442587794589006, 'support': 7945.0}, 'weighted avg': {'precision':
0.9451309444918609, 'recall': 0.944115796098175, 'f1-score':
0.9443713549203824, 'support': 7945.0}}
Results:

Hyperparamters: {'n_neighbors': 4, 'weights': 'distance', 'algorithm':
'kd_tree', 'leaf_size': 25, 'p': 1, 'n_jobs': -1}
Mean macro f1-score: 0.94
Mean deceased f1-score: 0.92
Mean overall accuracy: 0.94

2. Without Scalers(), with Hyperparam tuning():
Fitting 5 folds for each of 56 candidates, totalling 280 fits
K-Nearest best params: {'algorithm': 'kd_tree', 'leaf_size': 25,
'n_jobs': -1, 'n_neighbors': 4, 'p': 1, 'weights': 'distance'}
K-Nearest best score: 0.9363395125516248
K-Nearest Train Accuracy: 0.99
K-Nearest Test Accuracy: 0.94
K-Nearest Report:
{'deceased': {'precision': 0.9013878743608473, 'recall':
0.9409073579870377, 'f1-score': 0.9207237455698564, 'support':
2623.0}, 'hospitalized': {'precision': 0.9739096573208723, 'recall':
0.9388138138138138, 'f1-score': 0.956039755351682, 'support': 2664.0},
'non_hospitalized': {'precision': 0.9594543387646836, 'recall':
0.9525959367945824, 'f1-score': 0.9560128374551634, 'support':
2658.0}, 'accuracy': 0.944115796098175, 'macro avg': {'precision':
0.9449172901488011, 'recall': 0.9441057028651446, 'f1-score':
0.9442587794589006, 'support': 7945.0}, 'weighted avg': {'precision':
0.9451309444918609, 'recall': 0.944115796098175, 'f1-score':
0.9443713549203824, 'support': 7945.0}}

```

7. Overfitting

Random Forest AUC-ROC Plotting

```

def tune_hparam_auc_acc(model, x_train, x_test, y_train, y_test,
values, params, param: str):
    train_results = []
    test_results = []
    train_accuracies = []
    test_accuracies = []

    for value in values:
        params[param] = value
        train_acc, test_acc, _, _, avg_train_auc, avg_test_auc =
model(x_train, x_test, y_train, y_test, params)
        train_results.append(avg_train_auc)
        test_results.append(avg_test_auc)

```

```

        train_accuracies.append(train_acc)
        test_accuracies.append(test_acc)

    values = ['None' if v is None else v for v in values]
    plt.figure()
    line1, = plt.plot(values, train_results, 'b', label="Train AUC")
    line2, = plt.plot(values, test_results, 'r', label="Test AUC")
    line3, = plt.plot(values, train_accuracies, 'g', label="Train
Acc")
    line4, = plt.plot(values, test_accuracies, 'y', label="Test Acc")

    plt.legend(handler_map = {line1: HandlerLine2D(numpoints=2)})
    plt.ylabel('AUC/Acc score')
    plt.xlabel(param)
    plt.title('Different Values of ' + param + ' on Accuracy and AUC')
    plt.show()

def plot_auc_rf_tuning(rfmodel, x_train, x_test, y_train, y_test):
    n_estimators = [1, 2, 4, 8, 16, 32, 64, 128] # >40 is good
    tune_hparam_auc_acc(model_rf, x_train, x_test, y_train, y_test,
                        n_estimators, {'random_state': RANDOM_STATE},
                        'n_estimators')
    max_depths = range(1, 32) # >20 is good
    tune_hparam_auc_acc(model_rf, x_train, x_test, y_train, y_test,
                        max_depths, {'n_estimators': 40, 'random_state':
RANDOM_STATE}, 'max_depth')
    max_features = ['sqrt', 'log2', len(x_train[0]), None] # sqrt is
best
    tune_hparam_auc_acc(model_rf, x_train, x_test, y_train, y_test,
                        max_features, {'n_estimators': 40,
'random_state': RANDOM_STATE}, 'max_features')
    # 1% to 50% of total data
    min_samples_splits = np.linspace(0.01, 0.5, 50, endpoint=True) #
Lower is better (Only raise if not enough time/resources)
    tune_hparam_auc_acc(model_rf, x_train, x_test, y_train, y_test,
                        min_samples_splits, {'n_estimators': 40,
'random_state': RANDOM_STATE}, 'min_samples_split')
    # 1% to 50% of total data
    min_samples_leafs = np.linspace(0.01, 0.5, 50, endpoint=True) #
Lower is better (Only raise if not enough time/resources)
    tune_hparam_auc_acc(model_rf, x_train, x_test, y_train, y_test,
                        min_samples_leafs, {'n_estimators': 40,
'random_state': RANDOM_STATE}, 'min_samples_leaf')
    criterions = ['gini', 'entropy', 'log_loss'] # Gini is best
    tune_hparam_auc_acc(model_rf, x_train, x_test, y_train, y_test,
                        criterions, {'n_estimators': 40, 'random_state':
RANDOM_STATE}, 'criterion')
    return

```

Validation Curve Plotting

```
def plot_knn_k_tuning(X, y):  
    param_range = range(1, 11)  
    # Calculate accuracy on training and test set using the  
    # gamma parameter with 5-fold cross validation  
    train_score, test_score = validation_curve(KNeighborsClassifier(),  
X, y, param_name = "n_neighbors",  
param_range =  
param_range, cv = 5, scoring = "accuracy")  
    plt.figure()  
    line1, = plt.plot(param_range, np.mean(train_score, axis = 1),  
'r', label = "Train Accuracy")  
    line2, = plt.plot(param_range, np.mean(test_score, axis = 1), 'b',  
label = "Cross Validation (Test) Accuracy")  
    plt.legend(handler_map = {line1: HandlerLine2D(numpoints=2)})  
    plt.ylabel('Accuracy')  
    plt.xlabel('K Value')  
    plt.title('Validation Curve of KNN Classifier')  
    plt.show()  
  
def plot_rf_estimators_tuning(X, y):  
    param_range = range(0, 120, 20)  
    # Calculate accuracy on training and test set using the  
    # gamma parameter with 5-fold cross validation  
    train_score, test_score =  
validation_curve(RandomForestClassifier(), X, y, param_name =  
"n_estimators",  
param_range =  
param_range, cv = 5, scoring = "accuracy")  
    plt.figure()  
    line1, = plt.plot(param_range, np.mean(train_score, axis = 1),  
'r', label = "Train Accuracy")  
    line2, = plt.plot(param_range, np.mean(test_score, axis = 1), 'b',  
label = "Cross Validation (Test) Accuracy")  
    plt.legend(handler_map = {line1: HandlerLine2D(numpoints=2)})  
    plt.ylabel('Accuracy')  
    plt.xlabel('n_estimators')  
    plt.title('Validation Curve of RF Classifier')  
    plt.show()  
  
def plot_XGB_depth_tuning(X, y):  
    param_range = range(1, 16)  
    # Calculate accuracy on training and test set using the  
    # gamma parameter with 5-fold cross validation  
    train_score, test_score =
```

```

validation_curve(xgb.XGBClassifier(objective="multi:softprob",
random_state=RANDOM_STATE),

X, y, param_name =
"max_depth",
param_range =
param_range, cv = 5, scoring = "accuracy")
plt.figure()
line1, = plt.plot(param_range, np.mean(train_score, axis = 1),
'r', label = "Train Accuracy")
line2, = plt.plot(param_range, np.mean(test_score, axis = 1), 'b',
label = "Cross Validation (Test) Accuracy")
plt.legend(handler_map = {line1: HandlerLine2D(numpoints=2)})
plt.ylabel('Accuracy')
plt.xlabel('max_depth')
plt.title('Validation Curve of XGBoost Classifier')
plt.show()

def plot_XGB_estimators_tuning(X, y):

    param_range = range(100, 900, 100)
    # Calculate accuracy on training and test set using the
    # gamma parameter with 5-fold cross validation
    train_score, test_score = validation_curve(xgb.XGBClassifier

(objective="multi:softprob", random_state=RANDOM_STATE,
colsample_bytree = 0.1),

X, y, param_name =
"n_estimators",
param_range =
param_range, cv = 5, scoring = "accuracy")
plt.figure()
line1, = plt.plot(param_range, np.mean(train_score, axis = 1),
'r', label = "Train Accuracy")
line2, = plt.plot(param_range, np.mean(test_score, axis = 1), 'b',
label = "Cross Validation (Test) Accuracy")
plt.legend(handler_map = {line1: HandlerLine2D(numpoints=2)})
plt.ylabel('Accuracy')
plt.xlabel('n_estimators')
plt.title('Validation Curve of XGBoost Classifier')
plt.show()


```

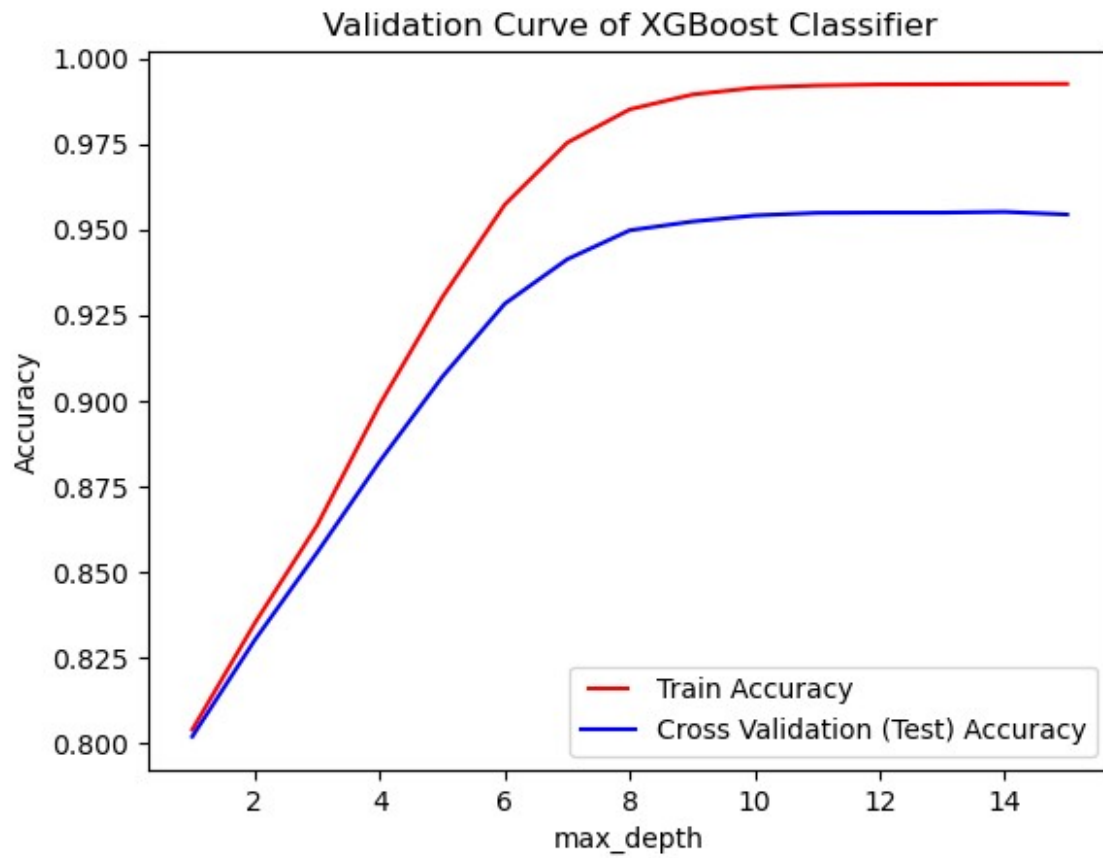
Driver Code

```


plot_knn_k_tuning(X, y)
plot_XGB_depth_tuning(X, y)
plot_XGB_estimators_tuning(X, y)
plot_rf_estimators_tuning(X, y)
plot_auc_rf_tuning(model_rf, x_train, x_test, y_train, y_test)

```

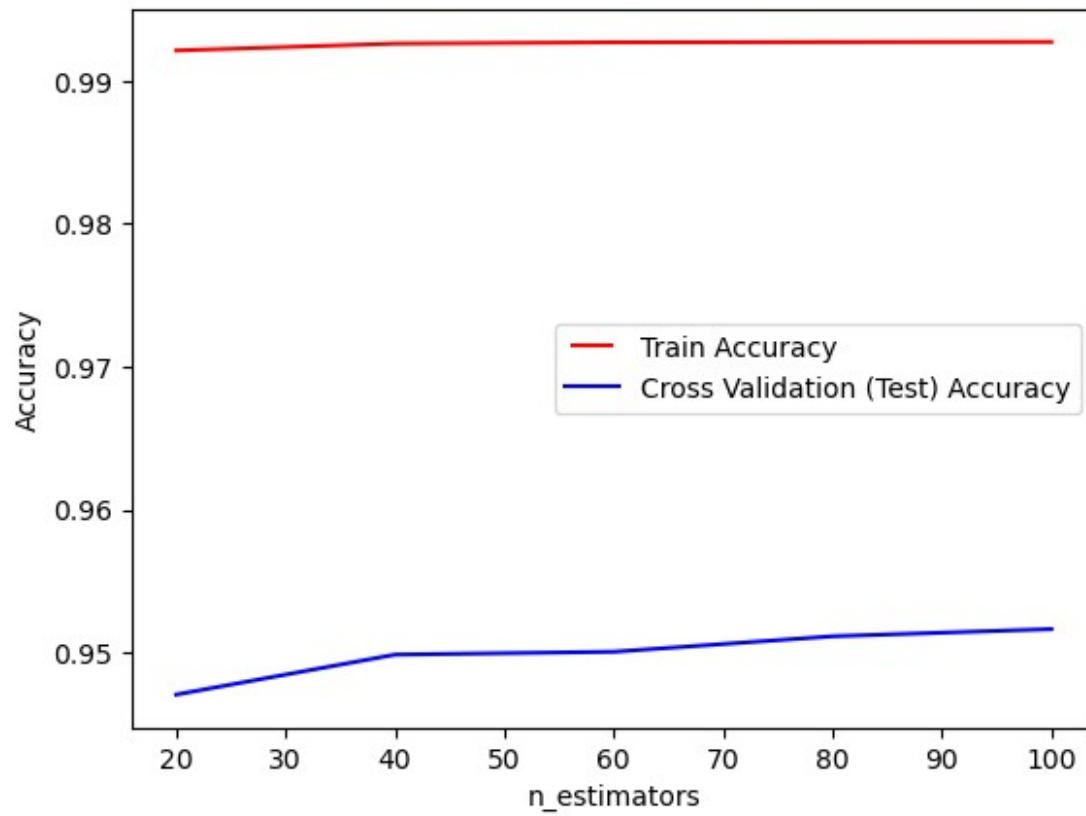





Validation Curve of XGBoost Classifier





Validation Curve of RF Classifier






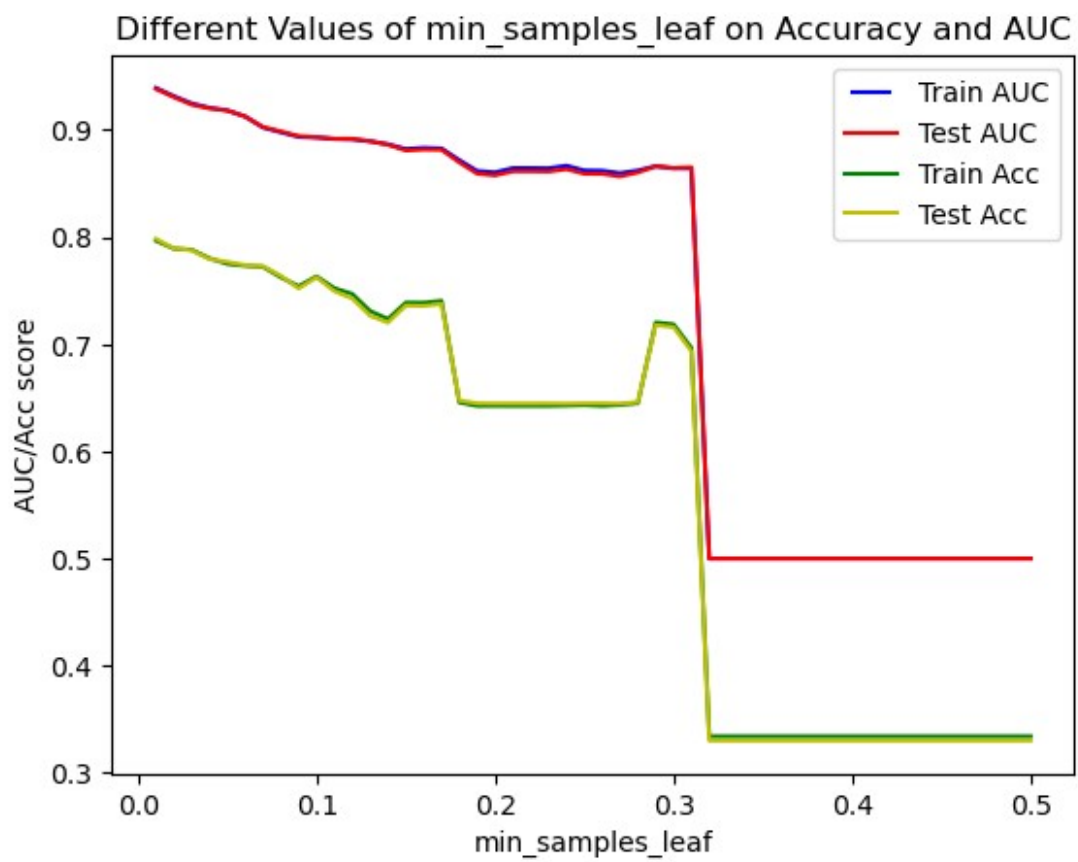
Different Values of max_depth on Accuracy and AUC

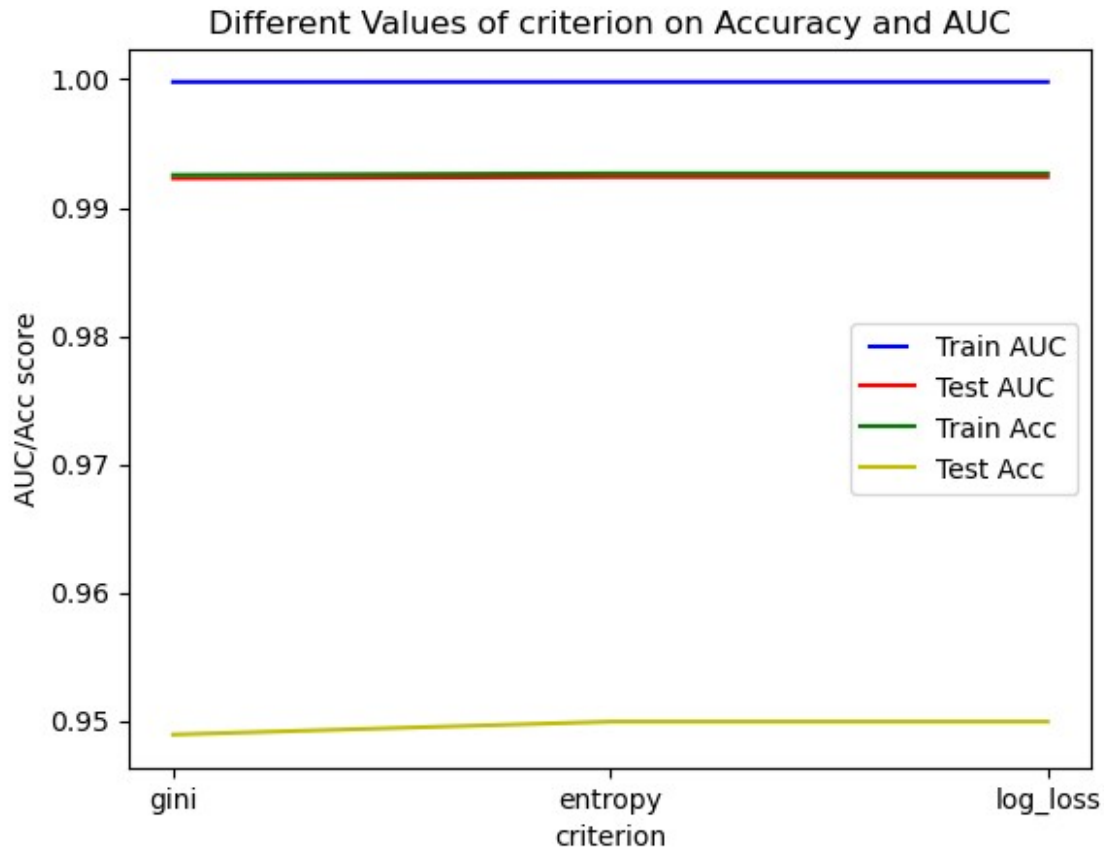




Different Values of min_samples_split on Accuracy and AUC







9. Prediction on the Test Set

Results

```
'''For labelling test data without predictions of the outcome_group'''
def create_submission_file(y_preds, file_name):
    with open(file_name, "w") as csvfile:
        wr = csv.writer(csvfile, quoting=csv.QUOTE_ALL)
        wr.writerow(["Id", "Prediction"])
        for i, pred in enumerate(y_preds):
            wr.writerow([str(i), str(pred)])
    return

clf = xgb.XGBClassifier(objective="multi:softprob",
                        random_state=RANDOM_STATE, **xgb_params)    # for multi-class
                             classification
clf.fit(x_train, y_train)
predictions = clf.predict(X_test_cases)
create_submission_file(y_preds=predictions, file_name="submission_
%s.csv" % "xgb")

clf = RandomForestClassifier(**rf_params)    # for multi-class
                             classification
```

```
clf.fit(x_train, y_train)
predictions = clf.predict(X_test_cases)
create_submission_file(y_preds=predictions, file_name="submission_
%s.csv" % "rf")

clf = KNeighborsClassifier(**knn_params)      # for multi-class
classification
clf.fit(x_train, y_train)
predictions = clf.predict(X_test_cases)
create_submission_file(y_preds=predictions, file_name="submission_
%s.csv" % "knn")
```